

Random effects

Olav N. R. Breivik

Thanks to Anders Nielsen for borrowing course material




Norsk Regnesentral
ANVENDT DATAFORSKNING

The Normal Distribution

- A continuous probability distribution on $(-\infty, \infty)$
- Probability density function:

$$\phi(x|\mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- Mean is μ and standard deviation is σ
- The interval $(\mu - 2\sigma, \mu + 2\sigma)$ contains 95%



Multivariate normal distribution

- In the multivariate case we write $x \sim N(\mu, \Sigma)$
- The density of x is

$$\phi(x|\mu, \Sigma) = \frac{1}{(2\pi)^{k/2} \sqrt{|\Sigma|}} \exp\left(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu)\right)$$

Example: Paired observations

- Two methods A and B for measure blood cell count (to check for doping)
- Paired study, want to see if both have same expectation

Person ID	Method A	Method B
1	5.5	5.4
2	4.4	4.9
3	4.6	4.5
4	5.4	4.9
5	7.6	7.2
6	5.9	5.5
7	6.1	6.1
8	7.8	7.5
9	6.7	6.3
10	4.7	4.2

- It should be expected that measures from same person are correlated
- A paired t-test gives p-value of 5.1%, which is a boarder line result.
- But more data is available
 - However with only one method for each person

In addition 10 persons were measured with only one method

- Want to use all data
- Not possible to use t-test with all data
- But it is possible by using latent effects!
- Consider the following model:

$$C_i = \alpha(M_i) + B(P_i) + \epsilon_i \text{ for measure } i = 1, \dots, 30$$

where $\alpha(M_i)$ is the fixed method effect

$B(P_i) \sim N(0, \sigma_p^2)$ is a random person effect (20 person effects in total)

$\epsilon_i \sim N(0, \sigma_\epsilon^2)$ is independent measurement noise

- This model uses all data and gives a 95% C.I. for method bias $\alpha(A) - \alpha(B)$ equal (0.04, 0.41)
- Notice the slightly significant method bias

Person ID	Method A	Method B
1	5.5	5.4
2	4.4	4.9
3	4.6	4.5
4	5.4	4.9
5	7.6	7.2
6	5.9	5.5
7	6.1	6.1
8	7.8	7.5
9	6.7	6.3
10	4.7	4.2
11		5.1
12		4.4
13		4.5
14		5.3
15		7.5
16	5.7	
17	6.0	
18	7.5	
19	6.5	
20	4.2	

Random effect model

For *fixed effect models* we have

- Random variables we observe (the response)
- Model parameters we want to estimate

For *random effect models* we have

- Random variables we observe (the response)
- Random variables we do **NOT** observe
- Model parameters we want to estimate

This model class is very usefull and goes under many names:

- | | | |
|--------------------------|------------------------------|------------------|
| • state space models | • random effect models | • mixed models |
| • latent variable models | • latent hierarchical models | • frailty models |
| • hidden Markov models | | |


Latent random effects

- Assume we have:

Observations: y

NOT observed random effects: u

Parameters (θ) in the model: $(y, u) \sim D(\theta)$



- How do we estimate our parameters when some of our random variables are not observed?


Latent random effects

- 1 Joint model (banana) is determined by model parameters (θ)
- 2 Marginal model is calculated from joint by integration
- 3 Marginal is matched to data (what we observe)

The marginal likelihood is:

$$L_M(y, \theta) = \int L(y, u, \theta) du.$$

Match your data with the marginal likelihood




Examples of use

- Time series models via state space models
 - Correlation via hidden processes
- Spatial statistics
 - Correlation introduces with hidden field
- Something unobserved gives extra variation
 - Overdispersion in Poisson via negative binomial


Knowledge about the latent process gives you predictive power!

Latent process in time



Olav N. R. Breivik

Coverage issues




Random effects

Latent random effects

- The marginal likelihood is:

$$L_M(\theta, y) = \int_{\mathbb{R}^q} L(\theta, u, y) du$$

- θ is model parameters
 - y is the observed random values (the observations)
 - u is the NOT observed random values
- How to calculate the integral?
 - Numerical integration not practical (need a lot of integration points)
 - Seldom an analytical solution
 - Solution: Approximate with use of Taylor-approximation



Latent random effects

We need to approximate the difficult integral

$$L_M(\theta, y) = \int_{\mathbb{R}^q} L(\theta, u, y) du.$$

Solution:

- Let $\ell(\theta, u, y) = \log L(\theta, u, y)$. Note that

$$\ell(\theta, u, y) \approx \ell(\theta, \hat{u}_\theta, y) - \frac{1}{2}(u - \hat{u}_\theta)^t (-\ell''_{uu}|_{u=\hat{u}_\theta})(u - \hat{u}_\theta)$$

is a 2.order Taylor approximation around

$$\hat{u}_\theta = \operatorname{argmax}_u \ell(\theta, u, y)$$

for a given θ .

Thereby is:


$$\begin{aligned}
 L_M(\theta, y) &= \int_{\mathbb{R}^q} L(\theta, u, y) du \\
 &= \int_{\mathbb{R}^q} \exp \left(\ell(\theta, u, y) \right) du \\
 &\approx \int_{\mathbb{R}^q} \exp \left(\ell(\theta, \hat{u}_\theta, y) - \frac{1}{2} (u - \hat{u}_\theta)^t (-\ell''_{uu}|_{u=\hat{u}_\theta}) (u - \hat{u}_\theta) \right) du \\
 &= L(\theta, \hat{u}_\theta, y) \int_{\mathbb{R}^q} \exp \left(-\frac{1}{2} (u - \hat{u}_\theta)^t (-\ell''_{uu}|_{u=\hat{u}_\theta}) (u - \hat{u}_\theta) \right) du \\
 &= L(\theta, \hat{u}_\theta, y) \sqrt{\frac{(2\pi)^q}{|-\ell''_{uu}|_{u=\hat{u}_\theta}|}}
 \end{aligned}$$

The last step is obtained by observing that the integrand has a Gaussian shape.

Taking the logarithm gives:

$$\ell_M(\theta, y) \approx \ell(\theta, \hat{u}_\theta, y) - \frac{1}{2} \log(|-\ell''_{uu}|_{u=\hat{u}_\theta}|) + \frac{q}{2} \log(2\pi)$$

- This is the Laplace approximation
- **Mini exercises 1:**
 - Where do we utilize automatic differentiation here?
 - Why care about conditional independence?



Syntax

- Very simple syntax:

```
1 obj <- MakeADFun(nll, par, random = c("parName1", "parName2"))
```

- The parameters *parName1* and *parName2* are now integrated over with the Laplace approximation
- Just as before:

```
1 obj$fn()#Marginal likelihood  
2 obj$gn()#Gradient of marginal likelihood
```

Gradient based search

Fixed effect model:

```
opt = nlminb(obj$par,obj$fn,obj$gr, control = list(trace = 1))
outer mgc: 103.5774
0: 398.42982: 0.00000 0.00000
outer mgc: 112.5499
1: 281.66042: -0.465673 0.884957
outer mgc: 68.03514
2: 205.84484: -0.622067 2.87883
outer mgc: 45.85355
3: 202.84258: -0.821876 2.87009
outer mgc: 1.537242
4: 200.74572: -0.753957 2.68198
outer mgc: 0.6074853
5: 200.68054: -0.753045 2.55830
outer mgc: 0.07071213
6: 200.67113: -0.752270 2.59332
outer mgc: 0.002883933
7: 200.67099: -0.752283 2.58967
outer mgc: 8.665257e-05
8: 200.67099: -0.752282 2.58952
outer mgc: 5.93635e-06
9: 200.67099: -0.752282 2.58952
```

- mgc is maximum gradient component

Gradient based search

Random effect model:

```
> opt = nlminb(obj$par, obj$fn, obj$gr, control = list(trace = 1))
...
1:      867.24500: 0.132409 0.991195
iter: 1  value: 848.5753 mgc: 7.724378 ustep: 1
iter: 2  value: 847.7413 mgc: 0.5669531 ustep: 1
iter: 3  value: 847.741 mgc: 0.0157684 ustep: 1
iter: 4  value: 847.741 mgc: 1.253356e-05 ustep: 1
iter: 5  mgc: 8.010481e-12
iter: 1  mgc: 8.010481e-12
outer mgc: 100.4661
2:      847.59842: -0.726562 1.50322
iter: 1  value: 865.6699 mgc: 1.774005 ustep: 1
iter: 2  value: 865.4707 mgc: 0.1210295 ustep: 1
iter: 3  value: 865.4705 mgc: 0.002688139 ustep: 1
iter: 4  value: 865.4705 mgc: 4.835442e-06 ustep: 1
iter: 5  mgc: 1.364867e-11
iter: 1  mgc: 1.364867e-11
outer mgc: 25.57054
3:      816.46720: -0.521138 2.48189
```

- `mgc` stands for maximum gradient component
- Inner optimization:
 - Find maximum a posteriori estimates of latent effects
 - Marginalize over the latent effects using the Laplace approximation

Markov property

- We will use Gaussian Markov random fields
- Remember this important result:

Let \mathbf{Q} be the precision matrix of a Gaussian random field γ , then

$$Q_{i,j} = 0 \Leftrightarrow \gamma_i \text{ and } \gamma_j \text{ are conditionally independent.}$$

- A Gaussian Markov random field has a sparse \mathbf{Q} .
- Using sparse structures is essential for fast inference
- RTMB automatically detects and use sparse structures
 - This makes RTMB much faster than ADMB

Exercise


Assume we observe y_1, \dots, y_n where $y_i \sim N(\mu_i, \sigma_y)$ and

$$\log \mu_i = \beta_0 + \gamma_i$$

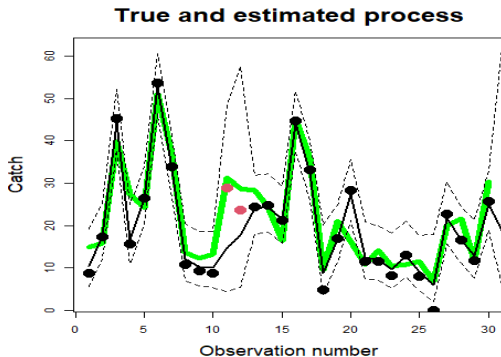
$$\gamma \sim N(0, \Sigma_{AR1}).$$

Here, γ is a stationary mean zero AR(1) process.

- We implemented this model with observed γ earlier today
- **Exercise 3a:** Implement and estimate the model with RTMB
- Data and code to get you started is provided in `ar1Latent.R`



- Missing observations, no problem
 - Use the data we have
 - Illustrate if time
- Can predict into the future




AR1: A stationary mean zero AR1 process has covariance matrix

$$\Sigma_{\text{AR1}} = \frac{\sigma^2}{1 - \rho^2} \begin{pmatrix} 1 & \rho & \rho^2 & \dots & \rho^{n-1} \\ \rho & 1 & \rho & \dots & \rho^{n-2} \\ \rho^2 & \rho & 1 & \dots & \rho^{n-3} \\ \vdots & \vdots & \vdots & \dots & \vdots \\ \rho^{n-1} & \rho^{n-2} & \rho^{n-3} & \dots & 1 \end{pmatrix}$$

- Coefficient of correlation is $\text{cor}(\gamma_i, \gamma_j) = \rho^{|j-i|}$
- The precision matrix is:

$$\Sigma_{\text{AR1}}^{-1} = \frac{1}{\sigma^2} \begin{pmatrix} 1 & -\rho & 0 & 0 & \dots & 0 & 0 & 0 \\ -\rho & 1 + \rho^2 & -\rho & 0 & \dots & 0 & 0 & 0 \\ 0 & -\rho & 1 + \rho^2 & -\rho & \dots & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \dots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \dots & -\rho & 1 + \rho^2 & -\rho \\ 0 & 0 & 0 & 0 & \dots & 0 & -\rho & 1 \end{pmatrix}$$

Obtained Sparsity structure in γ in Exercise 3



- Marginal log-likelihood:

$$\ell_M(\theta, y) \approx \ell(\theta, \hat{u}_\theta, y) - \frac{1}{2} \log(|-\ell''_{uu}|_{u=\hat{u}_\theta}|) + \frac{q}{2} \log(2\pi)$$

Summary


- A random effect is a random variable we do not observe
- We assign statistical structures to the random effects
- Marginalize over the latent variables
 - Laplace approximation

```
1 obj = MakeADFun(nll, par, random = c("nameOfLatentVariable1", ...))
```


- Optimize the marginal log-likelihood

```
1 opt = nlminb(obj$par, obj$fn, obj$gr)
```

Latent process in time



Spatial MAP



State-space stock assessment

- Remember we implemented a full parametric assessment model
- Let's go back to `introStockAssessment.pdf`
 - We will now move into a state-space assessment framework

Spatial modeling

- Next follow spatial modeling
- Will not go through spatial modeling in plenary
- Feel free to look at it and ask questions

Syntax for Markov models

- Let the random effect $\gamma \sim N(0, \mathbf{Q}^{-1})$

- Syntax in RTMB:

```
1 nll = nll - dgmrf(gamma, 0, Q, TRUE)
```

- Let the random effect $\gamma \sim N(0, \mathbf{Q}_{AR1}^{-1})$

- Syntax in RTMB:

```
1 nll = nll - dautoReg(gamma, 0, phi, log = TRUE, scale = sd, log = TRUE)
```

- Let the random effect $\gamma \sim N(0, \mathbf{Q}^{-1} \otimes \mathbf{Q}_{AR1}^{-1})$

- Syntax in RTMB:

```
1 f1 = function(x) dgmrf(gamma, 0, Q, TRUE)
2 f2 = function(x) dautoReg(x, phi=phi, log=TRUE)
3 nll = nll - dseparable(f1, f2)(gamma)
```

- Much used in spatio-temporal statistics:

$$\text{cov}(\gamma(s_1, t_1), \gamma(s_2, t_2)) = C^{(s)}(s_1, s_2) \cdot C^{(t)}(t_1, t_2)$$

Matern covariance structure with smoothness parameter equal 1

- Much used in spatial statistics
- Covariance is given by:

$$\text{Cov}(\gamma(s_1), \gamma(s_2)) = \sigma^2 \kappa \|s_1 - s_2\| K(\kappa \|s_1 - s_2\|).$$

- σ is the marginal variance
- κ is a spatial scale parameter
 - This variable can be increased on land to include spatial barriers
- Can be represented with a sparse precision matrix!
 - Use functionality from INLA/fmesher to extract Q and call

```
1 dgmrf(gamma, 0, Q, log = TRUE) #Likelihood contribution of latent effects
```

Spatial example





Figure: Area of bubbles is proportional to catch of haddock, crosses are zero-observations.

Spatial modeling with the SPDE-procedure

- Set up mesh

```
1 mesh = fmesher::fm_mesh_2d(locUTM,
2   max.edge = c(20, 50)) #inner and outer area
```



- Set up projection matrix

```
1 A = fmesher::fm_basis(mesh, loc = locUTM)
```



- Set up SPDE-matrices

```
1 spde = fmesher::fm_fem(mesh)
2 spdeMatrices = list(c0 = spde$c0, g1 = spde$g1, g2 = spde$g2)
```

- Calculate precision matrix

```
1 #Precision matrix Q, see Lindgren and Rue 2015 JSS p4
2 Q <- tau^2*(kappa^4*spdeMatrices$c0 + 2*kappa^2*
   spdeMatrices$g1 + spdeMatrices$g2)
```

$$Q = \frac{1}{\tau^2} (\kappa^4 C_0 + 2\kappa^2 G_1 + G_2)$$

Spatial modeling



The Institute of Marine Research conducts surveys to calculate fish stock. In this exercise, we use the catch of haddock from the Barents Sea winter survey.

Assume catch $C_i \sim \text{Pois}(\mu_i d_i)$, where

$$\log \mu(s_i) = \beta_0 + \delta(s_i).$$

Here, $\delta \sim N(0, \Sigma_{\text{Matern}})$ and d_i is distance trawled

- **Exercise 4 a)** Estimate the model
 - See `spde.R` to get you started
- **Exercise 4 b)** Estimate $\sum_{s \in \Omega} \mu(s)$
- Ω is a dense grid covering the survey area given in `predPoints.rds`



Simulate from the model

- Inform what are the observations

```
1 nll = function(par) {  
2   ...  
3   y = OBS(y)  
4   ...  
5 }
```

- Simulate from the model given hyperparameters

```
1 obj$simulate()
```